

```
In [1]: #importing necessary libraries
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings('ignore')
%matplotlib inline
```

```
In [2]: #reading the dataset from github
df = pd.read_csv('https://raw.githubusercontent.com/amberkakkar01/Prediction-of-Wine-Quality/master/winequality.csv')
```

```
In [3]: #dataframe head..giving first five rows of data
df.head()
```

Out[3]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol
0	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4
1	7.8	0.88	0.00	2.6	0.098	25.0	67.0	0.9968	3.20	0.68	9.8
2	7.8	0.76	0.04	2.3	0.092	15.0	54.0	0.9970	3.26	0.65	9.8
3	11.2	0.28	0.56	1.9	0.075	17.0	60.0	0.9980	3.16	0.58	9.8
4	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4

```
In [4]: #finding the columns
df.columns
```

```
Out[4]: Index(['fixed acidity', 'volatile acidity', 'citric acid', 'residual sugar',
               'chlorides', 'free sulfur dioxide', 'total sulfur dioxide', 'density',
               'pH', 'sulphates', 'alcohol', 'quality'],
              dtype='object')
```

Observations: Here we found 12 columns(Features), where 12 are independent features and 1 is dependent feature

In [5]: *#finding data type of the dataset*  
df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1599 entries, 0 to 1598
Data columns (total 12 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   fixed acidity                         1599 non-null   float64
1   volatile acidity                     1599 non-null   float64
2   citric acid                          1599 non-null   float64
3   residual sugar                       1599 non-null   float64
4   chlorides                           1599 non-null   float64
5   free sulfur dioxide                 1599 non-null   float64
6   total sulfur dioxide                 1599 non-null   float64
7   density                             1599 non-null   float64
8   pH                                  1599 non-null   float64
9   sulphates                           1599 non-null   float64
10  alcohol                             1599 non-null   float64
11  quality                             1599 non-null   int64
dtypes: float64(11), int64(1)
memory usage: 150.0 KB
```

Observations: We found 12 features data type is float and 1 dependent features'is integer

In [6]: *#checking null values of data*  
df.isnull().sum()

```
Out[6]: fixed acidity      0
volatile acidity      0
citric acid           0
residual sugar        0
chlorides             0
free sulfur dioxide    0
total sulfur dioxide   0
density               0
pH                   0
sulphates             0
alcohol               0
quality               0
dtype: int64
```

Observations: No null values found here

In [7]: *#Let us check the how many unique values are there in dependent feature*  
df.quality.unique()

```
Out[7]: array([5, 6, 7, 4, 8, 3], dtype=int64)
```

```
In [8]: #checking the values and counts of quality feature
df.quality.value_counts()
```

```
Out[8]: 5    681
        6    638
        7    199
        4     53
        8     18
        3     10
        Name: quality, dtype: int64
```

```
In [9]: #checking index index of quality feature
df.quality.value_counts().index
```

```
Out[9]: Int64Index([5, 6, 7, 4, 8, 3], dtype='int64')
```

```
In [10]: #checking index values of quality feature
df.quality.value_counts().values
```

```
Out[10]: array([681, 638, 199,  53,  18,  10], dtype=int64)
```

```
In [11]: #performing statistical analysis
df.describe(include='all').T
```

```
Out[11]:
```

	count	mean	std	min	25%	50%	75%	max
<b>fixed acidity</b>	1599.0	8.319637	1.741096	4.60000	7.1000	7.90000	9.200000	15.90000
<b>volatile acidity</b>	1599.0	0.527821	0.179060	0.12000	0.3900	0.52000	0.640000	1.58000
<b>citric acid</b>	1599.0	0.270976	0.194801	0.00000	0.0900	0.26000	0.420000	1.00000
<b>residual sugar</b>	1599.0	2.538806	1.409928	0.90000	1.9000	2.20000	2.600000	15.50000
<b>chlorides</b>	1599.0	0.087467	0.047065	0.01200	0.0700	0.07900	0.090000	0.61100
<b>free sulfur dioxide</b>	1599.0	15.874922	10.460157	1.00000	7.0000	14.00000	21.000000	72.00000
<b>total sulfur dioxide</b>	1599.0	46.467792	32.895324	6.00000	22.0000	38.00000	62.000000	289.00000
<b>density</b>	1599.0	0.996747	0.001887	0.99007	0.9956	0.99675	0.997835	1.00369
<b>pH</b>	1599.0	3.311113	0.154386	2.74000	3.2100	3.31000	3.400000	4.01000
<b>sulphates</b>	1599.0	0.658149	0.169507	0.33000	0.5500	0.62000	0.730000	2.00000
<b>alcohol</b>	1599.0	10.422983	1.065668	8.40000	9.5000	10.20000	11.100000	14.90000
<b>quality</b>	1599.0	5.636023	0.807569	3.00000	5.0000	6.00000	6.000000	8.00000

```
In [12]: x = df.iloc[:, :-1] #Seperating independent features from given dataset
```

```
In [13]: y = df.iloc[:, -1] #Seperating independent features from given dataset
```

In [14]: x

Out[14]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol
0	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56	9
1	7.8	0.880	0.00	2.6	0.098	25.0	67.0	0.99680	3.20	0.68	9
2	7.8	0.760	0.04	2.3	0.092	15.0	54.0	0.99700	3.26	0.65	9
3	11.2	0.280	0.56	1.9	0.075	17.0	60.0	0.99800	3.16	0.58	9
4	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56	9
...	...	...	...	...	...	...	...	...	...	...	...
1594	6.2	0.600	0.08	2.0	0.090	32.0	44.0	0.99490	3.45	0.58	10
1595	5.9	0.550	0.10	2.2	0.062	39.0	51.0	0.99512	3.52	0.76	11
1596	6.3	0.510	0.13	2.3	0.076	29.0	40.0	0.99574	3.42	0.75	11
1597	5.9	0.645	0.12	2.0	0.075	32.0	44.0	0.99547	3.57	0.71	10
1598	6.0	0.310	0.47	3.6	0.067	18.0	42.0	0.99549	3.39	0.66	11

1599 rows × 11 columns



In [15]: y

```
Out[15]: 0      5
1      5
2      5
3      6
4      5
..
1594   5
1595   6
1596   6
1597   5
1598   6
Name: quality, Length: 1599, dtype: int64
```

```
In [16]: #Performing standardization from sklearn
from sklearn.preprocessing import StandardScaler
```

```
In [17]: #splitting the data into train and test
from sklearn.model_selection import train_test_split
```

```
In [18]: x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.33, random_
```

In [19]: `x_train.head()`

Out[19]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol
<b>548</b>	12.4	0.350	0.49	2.6	0.079	27.0	69.0	0.99940	3.12	0.75	10
<b>355</b>	6.7	0.750	0.01	2.4	0.078	17.0	32.0	0.99550	3.55	0.61	12
<b>1296</b>	6.6	0.630	0.00	4.3	0.093	51.0	77.5	0.99558	3.20	0.45	9
<b>209</b>	11.0	0.300	0.58	2.1	0.054	7.0	19.0	0.99800	3.31	0.88	10
<b>140</b>	8.4	0.745	0.11	1.9	0.090	16.0	63.0	0.99650	3.19	0.82	9

In [20]: `x_test.head()`

Out[20]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol
<b>803</b>	7.7	0.56	0.08	2.50	0.114	14.0	46.0	0.9971	3.24	0.66	9
<b>124</b>	7.8	0.50	0.17	1.60	0.082	21.0	102.0	0.9960	3.39	0.48	9
<b>350</b>	10.7	0.67	0.22	2.70	0.107	17.0	34.0	1.0004	3.28	0.98	9
<b>682</b>	8.5	0.46	0.31	2.25	0.078	32.0	58.0	0.9980	3.33	0.54	9
<b>1326</b>	6.7	0.46	0.24	1.70	0.077	18.0	34.0	0.9948	3.39	0.60	10

In [21]: `scaler = StandardScaler()`

In [22]: `scaler`

Out[22]: `StandardScaler()`

In [23]: *#fitting the train data*  
`print(scaler.fit(x_train))`

`StandardScaler()`

In [24]: *#finding the mean*  
`print(scaler.mean_)`

```
[ 8.30345472  0.53246499  0.26933707  2.54691877  0.08772736 15.91223156
 46.76330532  0.99677933  3.31453782  0.65881419 10.41521942]
```

```
In [25]: #transform the training data  
scaler.transform(x_train)
```

```
Out[25]: array([[ 2.40069523, -1.03103722,  1.12742595, ..., -1.26096312,  
                  0.52726134, -0.01431863],  
               [-0.93967131,  1.22920403, -1.32502245, ...,  1.52622836,  
                  -0.28225704,  2.24363201],  
               [-0.99827424,  0.55113165, -1.37611513, ..., -0.74241587,  
                  -1.20742091, -0.86105011],  
               ...,  
               [-0.6466567 ,  0.49462562, -1.06955908, ...,  1.26695473,  
                  -0.68701624, -0.86105011],  
               [-0.23643625, -1.87862768,  0.4121285 , ...,  0.03540501,  
                  0.81637505,  1.39690052],  
               [-1.46709761, -1.3700734 , -0.04770558, ...,  0.48913386,  
                  -0.68701624,  2.90220094]])
```

```
In [26]: #another form of fit_transform  
x_train_tf = scaler.fit_transform(x_train)
```

```
In [27]: x_train_tf
```

```
Out[27]: array([[ 2.40069523, -1.03103722,  1.12742595, ..., -1.26096312,  
                  0.52726134, -0.01431863],  
               [-0.93967131,  1.22920403, -1.32502245, ...,  1.52622836,  
                  -0.28225704,  2.24363201],  
               [-0.99827424,  0.55113165, -1.37611513, ..., -0.74241587,  
                  -1.20742091, -0.86105011],  
               ...,  
               [-0.6466567 ,  0.49462562, -1.06955908, ...,  1.26695473,  
                  -0.68701624, -0.86105011],  
               [-0.23643625, -1.87862768,  0.4121285 , ...,  0.03540501,  
                  0.81637505,  1.39690052],  
               [-1.46709761, -1.3700734 , -0.04770558, ...,  0.48913386,  
                  -0.68701624,  2.90220094]])
```

## Model building

importing SVC model from sklearn

```
In [28]: from sklearn.svm import SVC
```

```
In [29]: model = SVC()
```

```
In [30]: #fitting our data into this model  
model.fit(x_train_tf, y_train )
```

```
Out[30]: SVC()
```

```
In [31]: model.score(x_train_tf, y_train) # checking training accuracy
```

```
Out[31]: 0.6778711484593838
```

Observation: Here we got 67% accuracy for train data

```
In [33]: #need to transform the test data  
x_test_tf = scaler.transform(x_test)
```

```
In [34]: x_test_tf
```

```
Out[34]: array([[ -3.53642095e-01,  1.55589436e-01, -9.67373729e-01, ...,  
                -4.83142240e-01,  6.85666499e-03, -7.66968836e-01],  
               [-2.95039173e-01, -1.83446751e-01, -5.07539654e-01, ...,  
                4.89133857e-01, -1.03395269e+00, -8.61050113e-01],  
               [ 1.40444556e+00,  7.77155778e-01, -2.52076279e-01, ...,  
                -2.23868614e-01,  1.85718440e+00, -4.84725007e-01],  
               ...,  
               [-2.02456406e-03, -1.25706134e+00,  6.16499196e-01, ...,  
                -2.94133945e-02,  6.42906824e-01,  1.96138818e+00],  
               [-6.06274859e-02,  4.50655383e+00, -1.37611513e+00, ...,  
                1.39659155e+00, -9.76129945e-01,  4.56087756e-01],  
               [ 4.66798811e-01,  7.20649747e-01, -6.09725004e-01, ...,  
                -2.23868614e-01, -6.87016236e-01, -7.66968836e-01]])
```

```
In [35]: y_predict = model.predict(x_test_tf)
```

```
In [36]: y_test
```

```
Out[36]: 803      6  
        124      5  
        350      6  
        682      5  
        1326     6  
        ..  
        813      4  
        377      7  
        898      7  
        126      5  
        819      5  
        Name: quality, Length: 528, dtype: int64
```

```
In [37]: y_predict
```

```
Out[37]: array([5, 5, 6, 5, 6, 5, 5, 5, 6, 6, 6, 5, 6, 5, 5, 7, 5, 6, 7, 5, 5, 5,
        6, 6, 5, 5, 6, 5, 5, 6, 5, 5, 6, 5, 6, 5, 6, 6, 5, 6, 5, 5, 6, 5,
        6, 6, 6, 6, 5, 6, 5, 5, 6, 7, 5, 5, 6, 5, 6, 5, 6, 6, 5, 5, 7, 5,
        6, 5, 7, 5, 6, 5, 6, 6, 6, 5, 7, 5, 6, 7, 5, 7, 5, 5, 6, 6, 5, 6,
        6, 5, 6, 5, 5, 6, 5, 6, 5, 6, 5, 5, 5, 5, 6, 6, 6, 6, 6, 5, 6, 5,
        6, 5, 6, 5, 6, 6, 6, 5, 5, 6, 6, 6, 6, 5, 5, 5, 6, 6, 5, 6, 6, 5,
        5, 6, 6, 5, 5, 5, 5, 6, 6, 6, 6, 5, 6, 5, 6, 5, 6, 5, 6, 6, 5, 6,
        6, 6, 5, 6, 5, 6, 7, 6, 6, 5, 5, 6, 5, 5, 5, 5, 5, 5, 6, 5, 7, 6,
        6, 5, 5, 5, 5, 7, 5, 7, 5, 6, 6, 6, 7, 5, 6, 6, 5, 6, 6, 5, 5, 5,
        6, 6, 5, 5, 5, 5, 7, 6, 5, 5, 6, 6, 7, 5, 6, 6, 6, 6, 6, 5, 6, 5,
        5, 6, 6, 6, 5, 5, 5, 7, 5, 5, 5, 5, 5, 6, 6, 5, 6, 5, 6, 6, 5, 5,
        6, 6, 5, 6, 6, 5, 6, 5, 6, 5, 5, 5, 5, 5, 5, 6, 6, 6, 6, 6, 5, 7,
        6, 7, 6, 5, 6, 6, 5, 6, 5, 5, 5, 5, 6, 6, 6, 5, 7, 5, 5, 5, 5, 6,
        5, 6, 5, 6, 5, 7, 6, 5, 5, 6, 5, 6, 6, 7, 5, 5, 6, 5, 5, 5, 6, 6,
        6, 7, 5, 5, 6, 5, 5, 6, 5, 5, 6, 5, 6, 5, 6, 5, 5, 5, 6, 5, 5, 6,
        6, 7, 5, 5, 6, 6, 6, 6, 5, 5, 6, 7, 5, 5, 6, 5, 6, 5, 6, 6, 6, 6,
        5, 5, 6, 6, 5, 5, 5, 5, 5, 5, 5, 6, 5, 6, 6, 5, 5, 5, 5, 5, 6, 6,
        5, 6, 5, 6, 5, 5, 5, 6, 6, 5, 6, 6, 6, 5, 5, 6, 5, 5, 5, 6, 6, 6,
        7, 6, 5, 6, 5, 5, 6, 5, 5, 6, 7, 6, 5, 5, 6, 7, 6, 6, 6, 6, 5, 7,
        5, 6, 6, 5, 5, 5, 6, 6, 5, 5, 6, 5, 7, 5, 5, 5, 6, 5, 5, 5, 5, 6,
        6, 6, 6, 5, 5, 5, 5, 6, 6, 5, 6, 6, 5, 5, 5, 6, 7, 6, 6, 5, 5, 5,
        5, 5, 6, 5, 5, 5, 5, 6, 7, 6, 6, 6, 5, 6, 6, 6, 6, 5, 6, 6, 6, 6,
        5, 6, 6, 6, 5, 5, 6, 6, 5, 5, 6, 5, 6, 5, 5, 5, 5, 5, 5, 5, 5, 5,
        6, 6, 6, 6, 6, 6, 5, 5, 5, 7, 6, 6, 6, 5, 5, 5, 6, 6, 7, 7, 5, 5],
      dtype=int64)
```

```
In [38]: # checking accuracy
from sklearn.metrics import accuracy_score
```

```
In [39]: accuracy_score(y_test,y_predict)
```

```
Out[39]: 0.5984848484848485
```

```
In [40]: # checking the accuracy Using Logistic regression model
from sklearn.linear_model import LogisticRegression
```

```
In [41]: model2 = LogisticRegression()
```

```
In [42]: model2
```

```
Out[42]: LogisticRegression()
```

```
In [43]: model2.fit(x_train_tf, y_train)
```

```
Out[43]: LogisticRegression()
```

```
In [44]: #predicting the test data in logistic regression model
y_predict2 = model2.predict(x_test_tf)
```



```
In [45]: y_predict2
```

```
Out[45]: array([5, 5, 6, 5, 6, 5, 5, 5, 6, 6, 6, 5, 6, 5, 5, 7, 5, 5, 7, 5, 5, 5,
        6, 6, 5, 5, 6, 5, 5, 6, 5, 5, 6, 5, 6, 5, 6, 6, 5, 6, 5, 5, 6, 5,
        6, 6, 6, 5, 5, 6, 5, 5, 6, 6, 5, 5, 6, 5, 6, 5, 5, 6, 5, 5, 7, 5,
        7, 5, 6, 5, 7, 5, 6, 6, 6, 5, 7, 6, 6, 7, 5, 7, 5, 6, 6, 6, 5, 6,
        6, 5, 6, 5, 6, 6, 5, 6, 5, 6, 5, 6, 5, 5, 6, 6, 6, 6, 5, 5, 6, 5,
        7, 5, 6, 5, 6, 6, 6, 5, 5, 6, 6, 5, 6, 5, 5, 5, 6, 6, 6, 6, 6, 5,
        5, 6, 6, 5, 5, 5, 5, 6, 6, 6, 6, 6, 6, 5, 6, 5, 6, 5, 6, 6, 5, 6,
        6, 6, 5, 6, 5, 6, 6, 6, 6, 5, 5, 6, 5, 5, 5, 5, 5, 5, 6, 5, 5, 6,
        6, 5, 5, 5, 5, 6, 5, 7, 5, 6, 6, 6, 7, 5, 6, 6, 6, 6, 6, 5, 5, 5,
        5, 6, 5, 5, 5, 5, 7, 6, 5, 5, 6, 6, 6, 5, 6, 6, 7, 6, 5, 5, 6, 5,
        5, 6, 6, 6, 5, 5, 5, 7, 5, 5, 5, 5, 6, 6, 6, 6, 5, 6, 5, 5, 5, 5,
        6, 6, 5, 5, 6, 5, 7, 5, 5, 6, 5, 5, 4, 5, 6, 6, 6, 7, 6, 6, 5, 7,
        6, 6, 5, 5, 6, 6, 5, 6, 5, 5, 5, 5, 6, 6, 6, 5, 7, 5, 5, 5, 5, 6,
        5, 6, 5, 6, 5, 7, 5, 5, 5, 6, 5, 6, 6, 7, 5, 5, 6, 5, 5, 5, 6, 6,
        6, 7, 6, 5, 5, 5, 5, 6, 5, 5, 6, 5, 6, 6, 6, 5, 5, 5, 6, 6, 5, 6,
        6, 7, 5, 5, 5, 6, 6, 7, 5, 5, 6, 7, 6, 5, 6, 5, 6, 5, 6, 6, 5, 7,
        5, 5, 6, 6, 5, 5, 5, 6, 6, 5, 6, 6, 5, 6, 6, 5, 5, 5, 5, 5, 6, 6,
        5, 6, 5, 6, 5, 5, 5, 6, 7, 5, 6, 6, 6, 5, 5, 6, 5, 6, 5, 5, 6, 6,
        6, 6, 5, 6, 5, 5, 6, 5, 5, 6, 7, 6, 6, 5, 5, 6, 6, 6, 6, 6, 5, 7,
        5, 6, 6, 5, 6, 5, 6, 6, 5, 5, 6, 5, 7, 5, 5, 5, 7, 5, 5, 5, 5, 6,
        6, 6, 7, 5, 5, 5, 5, 6, 6, 5, 6, 6, 6, 5, 5, 6, 7, 5, 6, 5, 5, 5,
        5, 5, 6, 5, 5, 5, 6, 7, 7, 6, 6, 6, 5, 5, 6, 6, 6, 6, 6, 6, 5, 6,
        5, 6, 6, 6, 5, 5, 6, 6, 5, 5, 6, 5, 6, 5, 5, 5, 5, 5, 5, 5, 6, 5,
        5, 6, 6, 6, 6, 6, 5, 3, 5, 7, 7, 5, 5, 5, 5, 5, 7, 6, 7, 7, 3, 5],
        dtype=int64)
```

```
In [46]: accuracy_score(y_test,y_predict2)
```

```
Out[46]: 0.571969696969697
```

Observations: We found that, we got only 57% accuracy using logistic Regression model

```
In [ ]:
```

```
In [ ]:
```